


# Introduction to Machine Learning with Python

## ML Fundamentals

*Dr. Süha Tuna*  
*İTÜ Informatics Institute*

# Before start

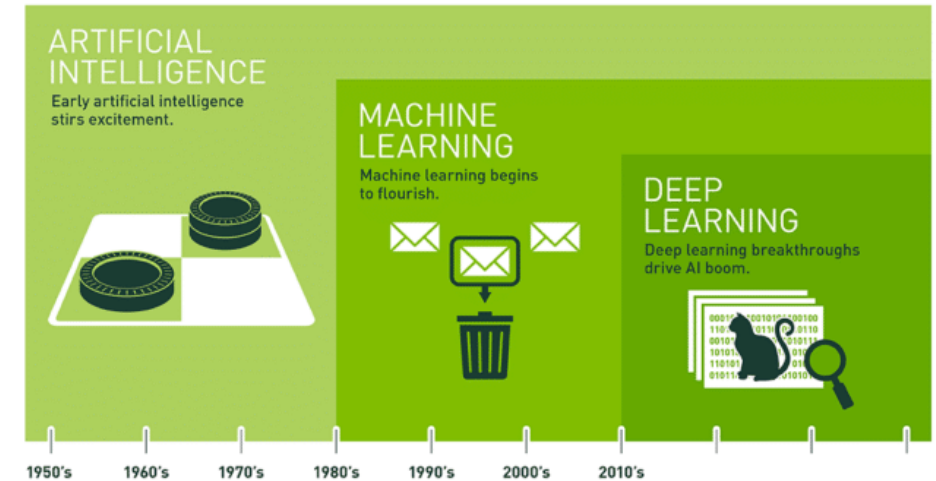
- Linear Algebra
  - Vectors, matrices
  - Eigenvalues, eigenvectors
  - Linear equation solvers
- Optimization
  - Gradients
  - Cost functions
- Statistics & Probability
  - Probability distributions
  - Likelihoods
  - Statistics



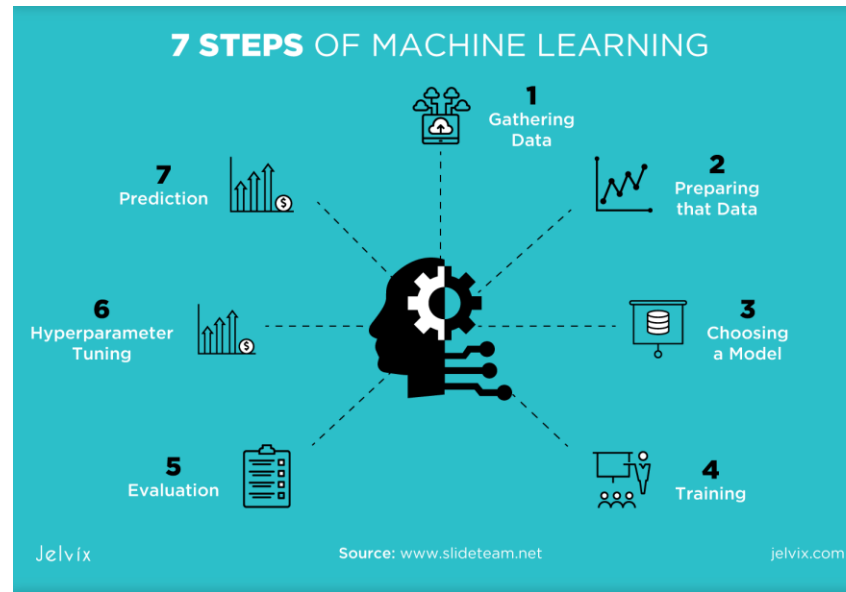
# What is ML?

## What is Machine Learning?

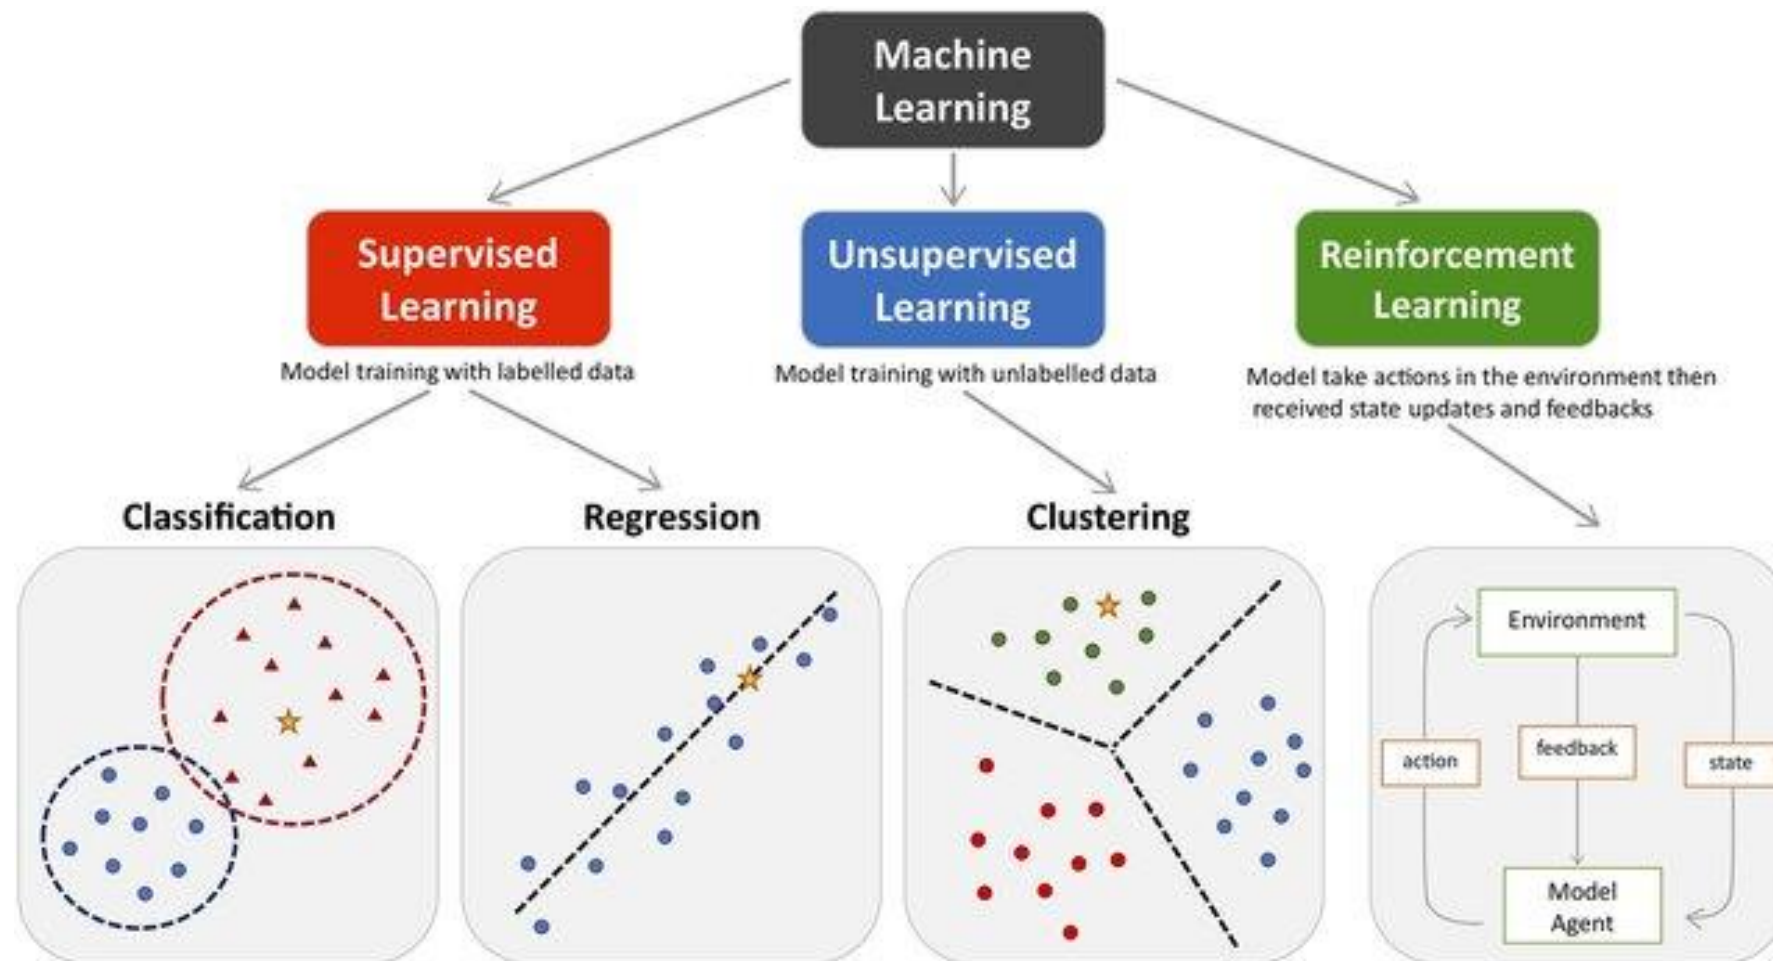
A branch of AI in which **machines analyze data, learn from it, identify patterns, and make decisions.**



Since an early flush of optimism in the 1950s, smaller subsets of artificial intelligence – first machine learning, then deep learning, a subset of machine learning – have created ever larger disruptions.



# ML Problem Types



# Bias-Variance Tradeoff

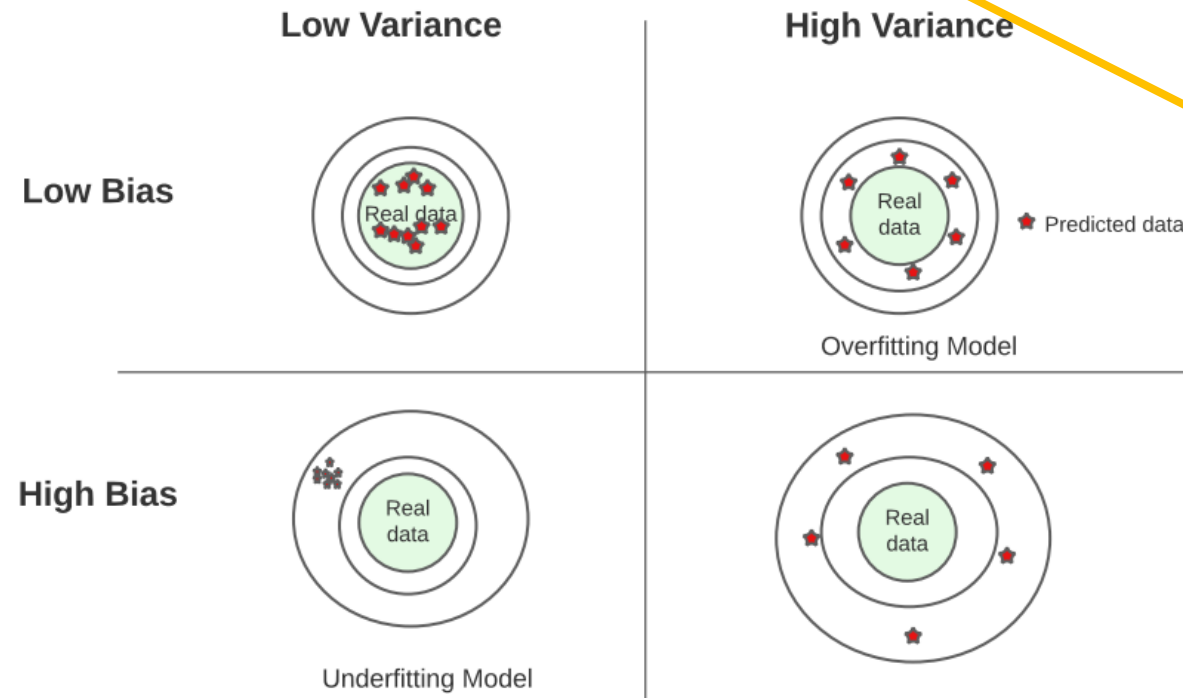
➤ A good model aims to **balance bias** and **variance** to **minimize total error**.

➤ **Total Error = Bias<sup>2</sup> + Variance + Irreducible Error**

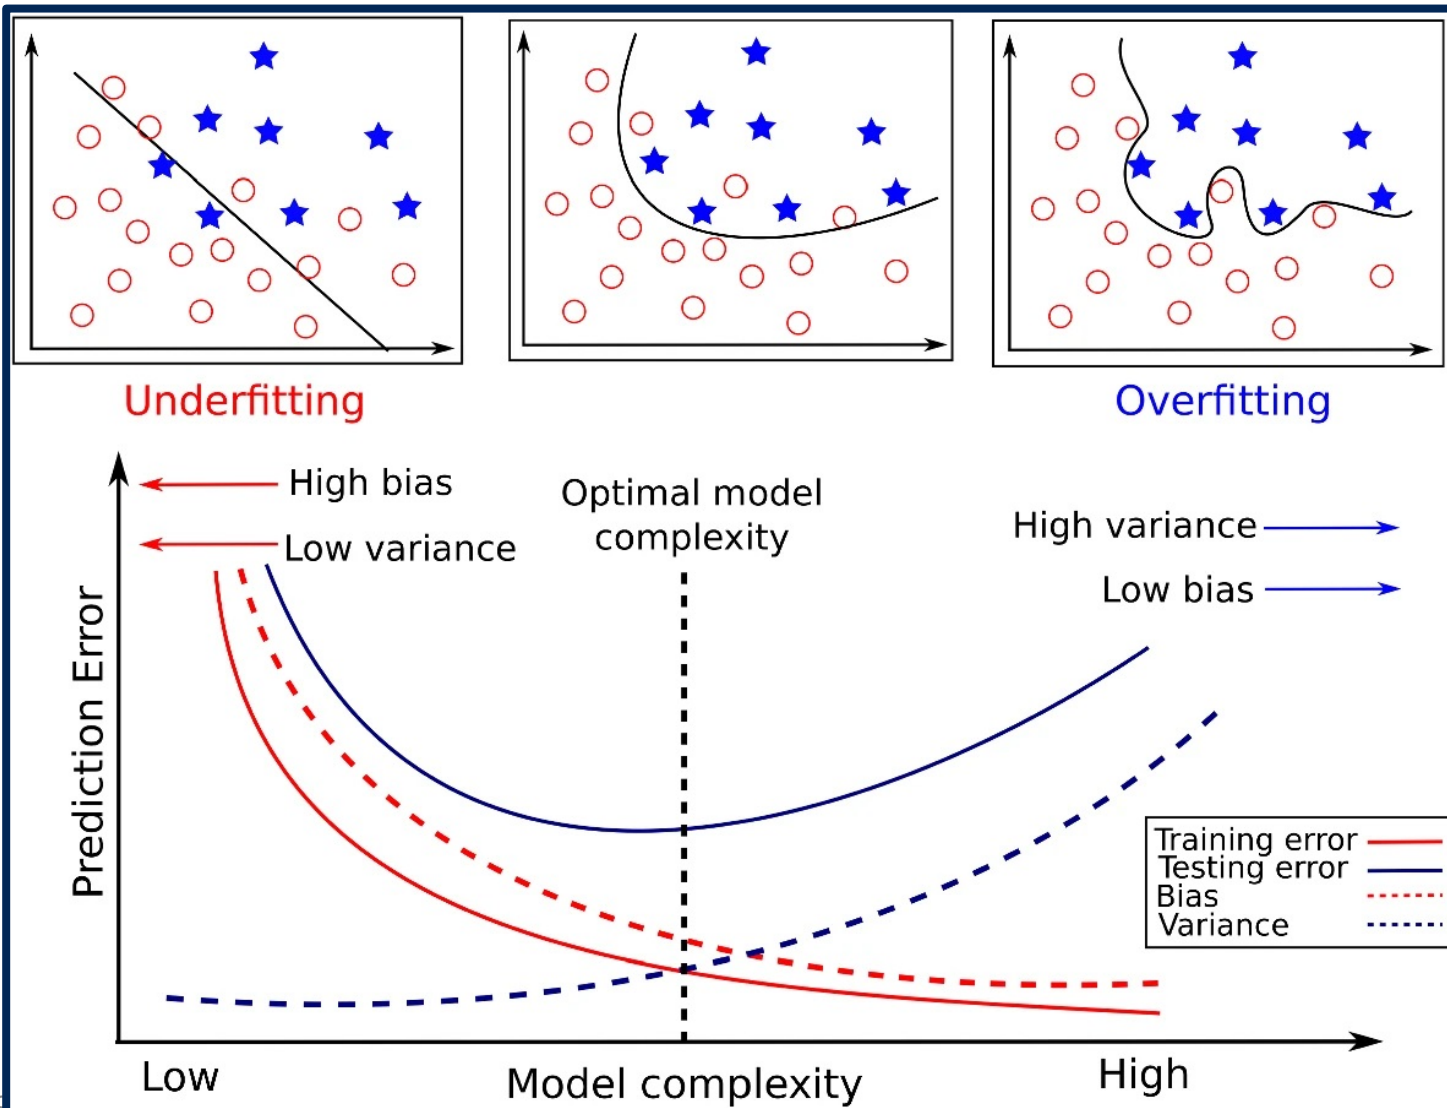
The error occurred from the model simplicity

The error occurred from the model complexity

Noise coming from the observations



# Bias-Variance Tradeoff



## ➤ Bias:

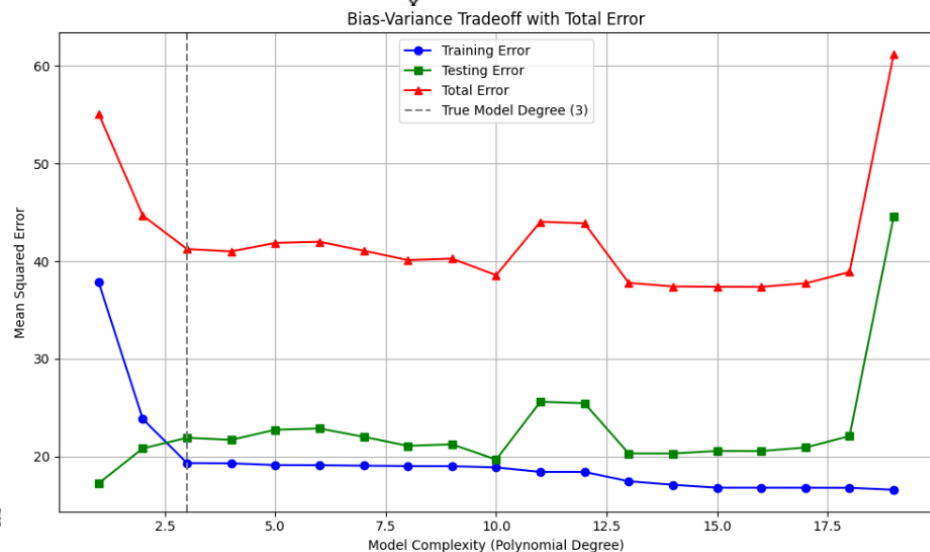
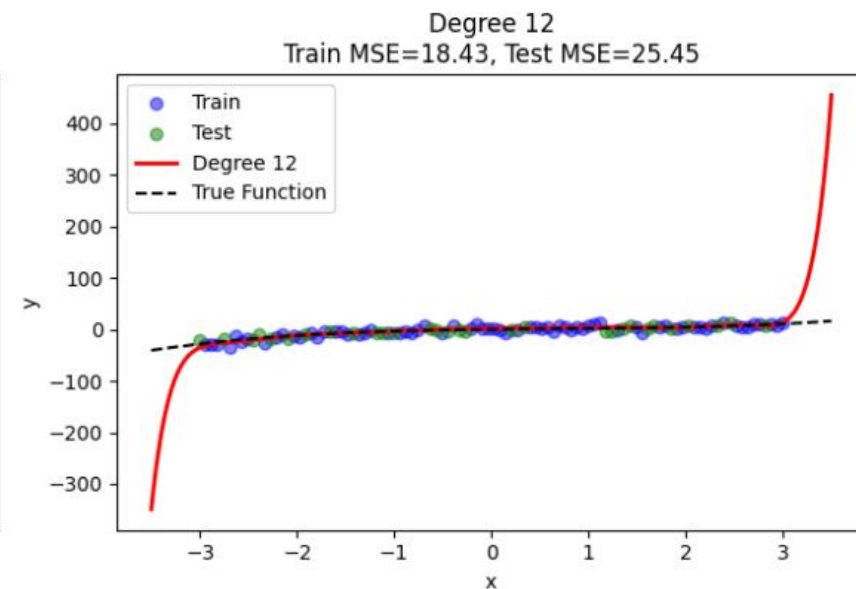

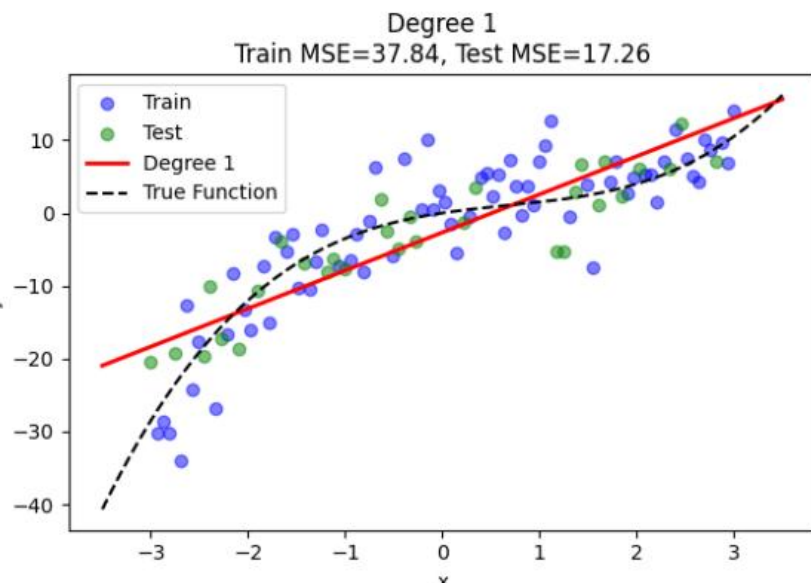
- The error due to **overly simplistic assumptions** in the learning algorithm.
- **High bias** means the **model is not complex enough** to capture the underlying patterns of the data.

## ➤ Variance:

- The error due to the **model's sensitivity to small fluctuations** in the training data.
- **High variance** means the **model fits the training data very closely but fails to generalize well**.



# Overfitting vs Ideal Fit vs Underfitting



- $y = 0.5x^3 - x^2 + 2x + \text{noise}$ 
  - Degree 1: Underfit (linear), high bias.
  - Degree 3: Ideal fit, low bias and variance.
  - Degree 12: Overfit, capturing noise, high variance.
- The error curve clearly drops then rises again after degree 3.

# Managing the tradeoff

## 1. Choose the Right Model Complexity

**Simple models** (e.g., linear regression) tend to have **high bias** and **low variance**, while **complex models** (e.g., deep neural networks) have **low bias** and **high variance**. Start with a simple model and increase complexity gradually, guided by validation performance.

## 2. Use More Training Data

**More training data** helps **reduce variance**, particularly for **complex models**, by enabling them to better capture the underlying data distribution.

## 3. Use Cross-Validation

Cross-validation (e.g., k-fold) provides a reliable estimate of model performance on unseen data. It helps in **identifying overfitting** (high variance) or **underfitting** (high bias).

## 4. Apply Regularization

Regularization techniques such as L2 (Ridge), L1 (Lasso), and dropout (for neural networks) penalize model complexity. They **reduce variance** by slightly increasing bias to prevent overfitting.



# Managing the tradeoff

## 5. Use Ensemble Methods

**Ensemble techniques** like bagging (e.g., Random Forest) reduce variance, while boosting (e.g., XGBoost) can **reduce bias**. They **combine** multiple **models** to improve **generalization**.

## 6. Feature Engineering and Selection

Creating informative features reduces bias, while **removing irrelevant or noisy features reduces variance**. Dimensionality reduction methods like PCA also help control variance.

## 7. Monitor Learning Curves

Plot training and validation errors against the number of training samples to diagnose bias and variance. High training and validation errors indicate high bias; low training error and high validation error indicate high variance.

# Pearson vs Spearman


© sunny kumar

Use Pearson correlation  
When there is linear Data



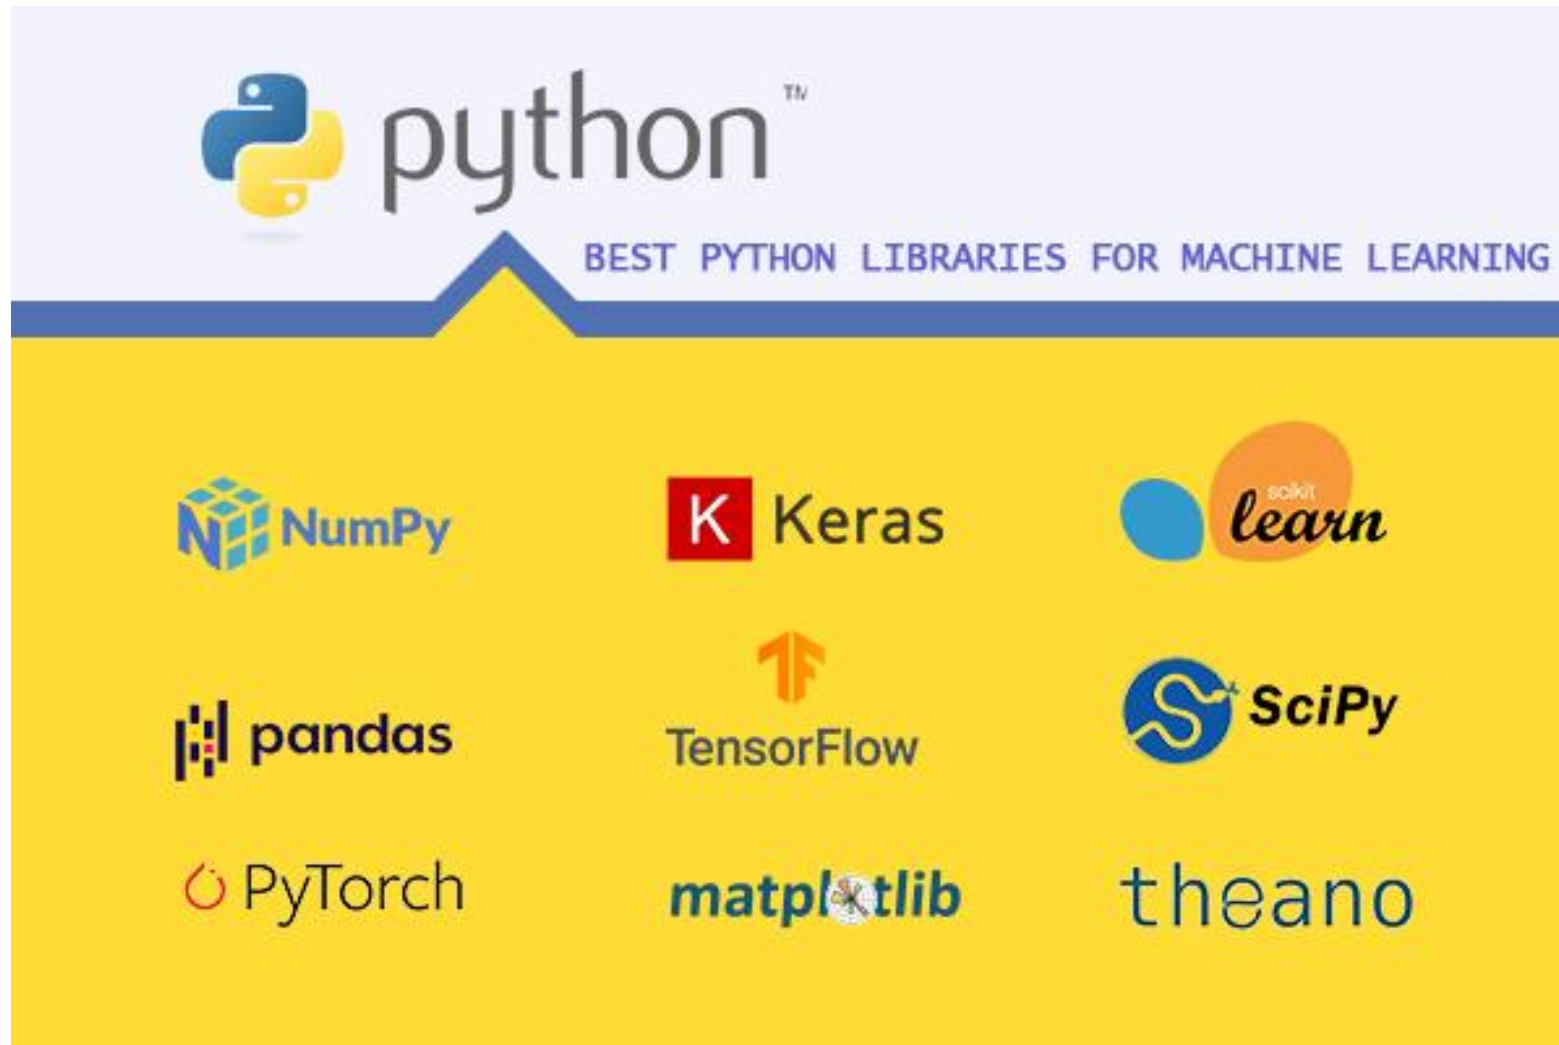
Pearson correlation is suitable when you want to measure the **strength and direction** of a linear relationship between two continuous variables

Use Spearman  
When there is non-linear Data



Spearman correlation is appropriate when you suspect a **monotonic**, but not necessarily linear, relationship between variables.

# ML Libraries in Python



# Python Libraries

## ➤ **NumPy** (Numerical Python)

- Purpose: Foundation for numerical computing in Python.
- Key Feature: ndarray, a powerful N-dimensional array object.
- Speed: Vectorized operations using compiled C backend (fast linear algebra).
- Uses: Matrix computations, random sampling, broadcasting, Fourier transforms.

## ➤ **Pandas**

- Purpose: Data manipulation & analysis with labeled axes.
- Core Structures: Series (1D) and DataFrame (2D).
- Features:
  - Intuitive slicing, filtering, merging, and groupby operations.
  - Handles missing data, time series, categorical variables.
- Best For: Tabular data processing and transformation pipelines.

# Python Libraries

## ➤ Scikit-learn

- Purpose: Standard ML toolkit for supervised & unsupervised learning.
- Strengths:
  - Unified API for models: `fit()`, `predict()`, `score()`
  - Pipelines, cross-validation, hyperparameter tuning (`GridSearchCV`, `RandomizedSearchCV`)
- Algorithms: SVM, Random Forest, KNN, PCA, KMeans, Logistic Regression, etc.
- Integration: Works seamlessly with NumPy/pandas.

## ➤ SciPy (Scientific Python)

- Purpose: Advanced scientific computing built on top of NumPy.
- Submodules:
  - `scipy.linalg`: Linear algebra (beyond NumPy)
  - `scipy.optimize`: Optimization & curve fitting
  - `scipy.spatial`: Distance metrics, KD-trees
  - `scipy.stats`: Statistical functions, distributions, hypothesis testing
- Use Case: Numerical routines required in engineering, physics, and ML.

# Plotting and Typical Workflow

## ➤ **Matplotlib**

- Purpose: 2D plotting library for visualizing data and models.
- Main Interface: pyplot (similar to MATLAB).
- Features: Line plots, scatter plots, bar charts, histograms, heatmaps. Highly customizable (styles, ticks, annotations).
- Extensions: Integrates with seaborn for statistical plots.

NumPy → pandas (data wrangling) → scikit-learn (ML modeling) → SciPy (scientific routines) →  
Matplotlib (visualization)



# Importing Libraries

```
import sys
print("Python version: {}".format(sys.version))

import pandas as pd
print("pandas version: {}".format(pd.__version__))

import matplotlib
print("matplotlib version: {}".format(matplotlib.__version__))

import numpy as np
print("NumPy version: {}".format(np.__version__))

import scipy as sp
print("SciPy version: {}".format(sp.__version__))

import IPython
print("IPython version: {}".format(IPython.__version__))

import sklearn
print("scikit-learn version: {}".format(sklearn.__version__))
```


# Supervised Learning Cases

Example	Task Type	Description
1 Spam Email Detection	Classification	Predict if an email is <b>spam or not</b> based on text features.
2 House Price Prediction	Regression	Predict the <b>price</b> of a house based on size, location, etc.
3 Medical Diagnosis	Classification	Classify whether a tumor is <b>benign or malignant</b> from imaging features.
4 Credit Scoring	Classification	Assess if a person is likely to <b>default on a loan</b> using credit history.
5 Stock Price Forecasting	Regression	Predict the <b>next day's stock price</b> using historical data.

# Unsupervised Learning Cases

Example	Task Type	Description
1 Customer Segmentation	Clustering	Group customers by purchasing behavior (e.g., into <b>segments</b> ).
2 Anomaly Detection	Outlier Detection	Detect <b>fraudulent transactions</b> without explicit labels.
3 Topic Modeling	Dimensionality Reduction	Discover <b>latent topics</b> in a collection of documents (e.g., LDA).
4 Image Compression	Feature Extraction	Reduce image size using <b>PCA</b> while preserving key information.
5 Gene Expression Analysis	Clustering	Cluster genes with similar expression patterns across samples.

# A First Application: Classifying Iris Species



## ➤ Classification

- Setosa
- Versicolor
- Virginica